

Design and Analysis of Algorithms

Module 4: Divide-and-Conquer

These notes are © Grant Williams, [CC BY-SA 4.0](#).

This is an open educational resource.

Feel free to submit fixes, improvements, and new material [here](#).

Introduction

In this class, we're going to learn an extremely common and important design technique for algorithms: divide and conquer.

We're also going to learn how to analyze these algorithms for their runtime bounds. Our traditional proof techniques will work, but there's a cool theorem that will analyze them for us much easier than an inductive proof.

Divide and conquer

Divide and conquer algorithms are algorithms that are broken up into three phases:

1. Divide: the data is broken into 2 or more pieces
2. Conquer: the problem is solved on at least one of the smaller pieces
3. Recombine (optional): the pieces are somehow merged back together

Sometimes this paradigm is called "divide-conquer-recombine", but the recombine part is normally implied.

We often abbreviate it "D&C".

An example

Rather than keep this theoretical, let's see an example of a common, useful D&C algo.

We'll prove that it's correct, prove its runtime bound the old fashioned way, and then see if we see any patterns as we explore other D&C algorithms.

The algorithm we look at first will be *Mergesort*.

However, to understand mergesort, we have to first understand merging.

Merge

Have you ever noticed that it's faster to sort two lists that are already sorted than one list that isn't sorted?

I notice this whenever I have to sort exams by student name. If I break the piles into pieces, maybe with a distribution sort, I can sort the small piles quickly, and then merge them all together.

But how does this work, and how fast is it?

The merge algorithm

Suppose we have two sorted lists: `a = [1, 2, 3, 4]` , `b = [0, 1, 2, 5]`

Each one is sorted, but they aren't sorted with respect to each other. So we can't just concatenate them.

Instead, there's a simple algorithm we can use to combine them. First, assume there is an array that is big enough to hold both of them:

```
buf = [_, _, _, _, _, _, _, _]
```

The merge algorithm (2)

Then, we need to track three variables: a separate index or pointer into `a`, `b`, and `buf`.

```
a = [1, 2, 3, 4], b = [0, 1, 2, 2], buf = [_, _, _, _, _, _, _, _]
```

^

^

^

Here, the carats represent the indices. For screen narration, let `i = 0` be the index into `a`, `j = 0` be the index into `b`, and `k = 0` be the index into `buf`.

The merge algorithm (3)

```
a = [1, 2, 3, 4], b = [0, 1, 2, 2], buf = [_, _, _, _, _, _, _, _]  
      ^               ^               ^, i = 0, j = 0, k = 0
```

First, we compare the elements at `a[i]` and `b[j]`.

We want to copy the smaller one into the buffer. Specifically, always take `a[i]` if `a[i] <= b[j]`, otherwise take `b[j]`. *

In this case, `b[j]` is strictly smaller, so we copy it into the buffer, and increment `j`, which stores how many elements from `b` have been taken. We also increment `k`, which stores how many elements have been written to the buffer.

```
a= [1, 2, 3, 4], b= [0, 1, 2, 2], buf = [0, _, _, _, _, _, _, _]  
      ^               ^               ^, i = 0, j = 1, k = 1
```

*: The reason it's important to prioritize `a` is for reasons of stability. We'll talk about stability and why it's useful later.

The merge algorithm (4)

```
a= [1, 2, 3, 4], b= [0, 1, 2, 2], buf = [0, _, _, _, _, _, _, _]  
      ^               ^               ^, i = 0, j = 1, k = 1
```

Now we compare `a[i] == 1` with `b[i] == 1`. These are equal, so we prefer to take from `a`. Again, this is for reasons of stability, which we'll talk about soon.

```
a= [1, 2, 3, 4], b= [0, 1, 2, 2], buf = [0, 1, _, _, _, _, _, _]  
      ^               ^               ^, i = 1, j = 1, k = 2
```

And then we take from `b`, because `1 < 2`:

```
a= [1, 2, 3, 4], b= [0, 1, 2, 2], buf = [0, 1, 1, _, _, _, _, _]  
      ^               ^               ^, i = 1, j = 2, k = 3
```

[What happens next?]

The merge algorithm (5)

take from `a` , because `2 <= 2` :

`a= [1, 2, 3, 4], b= [0, 1, 2, 2], buf = [0, 1, 1, 2, _, _, _, _]`

^

^

^

, `i = 2` , `j = 2` , `k = 4`

then, the next two take from `b` , because both elements are `2` , while `a[i]==3` :

`a= [1, 2, 3, 4], b= [0, 1, 2, 2], buf = [0, 1, 1, 2, 2, 2, _, _]`

^

~

^

, `i= 1` , `j= 4` , `k= 6`

At this point, `b` has no more elements (I marked the index with `~`). So we can safely copy the rest of `a` into `buf` at position `k` :

`buf = [0, 1, 1, 2, 2, 2, 3, 4]`

And now `buf` is sorted, and contains all the elements in both `a` and `b` .

Questions?

Merge's code

So, can we write merge in C?

I recommend doing it as a practice exercise later. I'm going to show you the code, but try doing it on your own without using the code as a reference.

Here's a simple implementation...

merge in C (scrunched to fit)

```
void merge(const int* restrict lo, const int* restrict hi,
           size_t n_lo, size_t n_hi, int* restrict merge_buf)
{
    int i_lo = 0, i_hi = 0, i_merge = 0;
    for (;;) {
        if (i_lo == n_lo) { // done taking from lo
            memcpy(merge_buf + i_merge, hi + i_hi, (n_hi - i_hi) * sizeof(int));
            break;
        } if (i_hi == n_hi) { // done taking from hi
            memcpy(merge_buf + i_merge, lo + i_lo, (n_lo - i_lo) * sizeof(int));
            break;
        }
        if (lo[i_lo] <= hi[i_hi]) { merge_buf[i_merge++] = lo[i_lo++]; }
        else { merge_buf[i_merge++] = hi[i_hi++]; }
    }
}
```

Understanding `merge`

First, let's look at the parameters:

- `lo` and `hi` are the two arrays we're merging. They're called that because they will come from the low and high halves of an array.
- `n_lo` and `n_hi` are the lengths of the two arrays. In practice, `n_lo - n_hi <= 1`, but we don't need that invariant for this function to work.
- Lastly, we have a pointer to the buffer that we're going to be merging in to. This is a `restrict` pointer, which you may not have seen before and which we'll talk about on the next slide.

Restrict

`restrict` means that this pointer is the *only* pointer that will be used to access the data it accesses. It is a promise to avoid aliasing, which is when two or more pointers point to the same thing.

If you say `int* restrict a = array1; int* restrict b = array2;`, you're saying that `a` and `b` will never point to the same array.

Why is this useful? Because if they could point to the same data, then many optimizations don't work.

Without `restrict` the compiler would have to consider the theoretical possibility that changing `merge_buf` could also change what is read from `lo` or `hi`, forcing it to reload the data after writing it.

Restrict (2)

The `restrict` keyword in C only applies to pointers, so it must go after the `*`.

Remember that `const int*` and `int const*` both mean a pointer that points to a constant int, so the pointer can change, but a new value cannot be written to the underlying memory. But `int* const` means that the pointer itself is constant (it cannot be moved to a different value) but the underlying value can be changed.

Restrict works the same way. It is an error to say `restrict int*` or `int restrict*`. It must come after the `*`.

Note: when referring to a pointer to an int that is constant, people who write `const int*` are called "left-consters", and people who write `int const*` are called "right-consters." The main benefit of being a right-conster is that the underlying types will line up if you have a non-const type a line below a const type or vice versa. The main drawback is that less experienced programmers will think that you're declaring a constant pointer. It's also kind of weird that a modifier is coming after the thing it modifies, which is why I usually don't do it.

Understanding merge (2)

After our parameters, we get to the declarations and loop.

The first thing we check for is whether `i_lo`, which stores how many things we have taken from the `lo` array, is equal to `n_lo`, which is the total number of things in `lo`.

```
int i_lo = 0, i_hi = 0, i_merge = 0;
for (;;) {
    if (i_lo == n_lo) { // done taking from lo
        memcpy(merge_buf + i_merge, hi + i_hi, (n_hi - i_hi) * sizeof(int));
        break;
    }
```

If it is, we just copy the rest of it into `merge_buf`.

We do the same thing with `i_hi`. If we've used up all the elements of `hi`, we copy everything remaining from `lo`.

Understanding merge (3)

Finally we reach the main logic:

```
if (lo[i_lo] <= hi[i_hi])  
    merge_buf[i_merge++] = lo[i_lo++];  
else  
    merge_buf[i_merge++] = hi[i_hi++];
```

If the next element of `lo` is less than or equal the next element of `hi`, we copy the next element of `lo` to the buffer. Otherwise we copy the next element of `hi`.

Eventually one of these arrays will be empty, and we'll transition to the first part: copying over the rest of the other array.

And then we're done.

Questions?

Proving it

How can we prove our merge is correct?

We need some kind of loop invariant. How about this one:

```
sorted(lo), sorted(hi), sorted(merge_buf[0 .. i_merge)),
```

```
all elements in merge buf  $\leq$  lo[i_lo] and also  $\leq$  hi[i_hi]
```

So, basically, both our sub-arrays are sorted, our merge buffer so far is sorted, and everything in the merge buffer is smaller than everything in either sub array.

Proving it (2)

Is this true when `lo` is empty? Then we need this to be true after:

```
sorted([]), sorted(hi), sorted(merge_buf[0 .. n))
```

```
everything in merge_buf <= lo (vacuous), everything in merge_buf <= hi
```

This checks out. `n = len(lo) + len(hi)`.

Once `lo` is empty, we have `sorted(merge_buf[0 .. len(lo) + i_hi])`

And this array is less than everything remaining in `hi`.

Then we concat the rest of that array to the end. Solid!

The same applies when `hi` is empty conversely.

Proving it (3)

But what if neither `lo` nor `hi` are empty?

Say that the front of `lo` is smaller or equal. We have that `everything in merge_buf <= lo[lo_i]`. So that means, `merge_buf ++ [head lo]` is sorted. And, we don't violate our `everything in merge_buf <= hi` assumption, because we took something that was smaller than everything in `hi`.

Likewise, if the front of `hi` is smaller, the invariant is maintained by the same logic.

Proving it (4)

We've shown that our loop invariant holds.

Once the loop finishes, `merge_i == n_lo + n_hi`, so the entire array is sorted.

Questions?

What about mergesort?

Merge is a useful algorithm, because it lets us, say, use different sorting algorithms and then stitch the results together.

But we can sort entirely by applying `merge` recursively.

Mergesort psuedocode

```
mergesort(arr) =  
    n = len(arr)  
    if n <= 1: return arr  
  
    first_hi = (n + 1) / 2    -- ceil of division  
    lo = arr[0 .. first_hi)  
    hi = arr[first_hi .. n)  
  
    lo = mergesort(lo)  
    hi = mergesort(hi)  
  
    return merge(lo, hi)
```

Understanding the psuedocode

First is our base case. Empty and singleton arrays are sorted, so just return.

Otherwise, we want to split the array into two pieces, its lower and upper halves.

If the array has an even length, we want both halves to be the same size. On the other hand, if it's odd, we want `lo` to be one element bigger. (Proving that $(n + 1) / 2$ gives us this is a practice exercise)

We then recursively apply `mergesort` on both halves.

Once it's done, we merge the results together.

Wait, really?

It seems strange to use recursion this way. We're doing recursion *before* the main logic

That recursion will keep recursing. It will keep cutting the lists into halves. Eventually, the halves will be of size 1. Then we hit the base case.

Once the base case returns, we merge together two size-1 arrays. Then that returns and we end up merging together two size-2 arrays. And so on


Tracing it

Consider `mergesort([5,3,1,2,4])`

First, it splits into `mergesort([5,3,1])` and `mergesort([2,4])`

- `mergesort([5, 3, 1])` splits into `mergesort([5, 3])` and `mergesort([1]) = [1]`
 - `mergesort([5, 3]) = mergesort([5]) = [5]` and `mergesort([3]) = [3]`.
 - We `merge([5], [3])` to get `[3, 5]`
 - we `merge([3, 5], [1])` to return `[1, 3, 5]`
- `mergesort([2, 4])` splits into `mergesort([2])` and `mergesort([4])`
 - We `merge([2], [4])` to return `[2, 4]`
- We `merge([1, 3, 5], [2, 4])` to get `[1, 2, 3, 4, 5]`

Visualizing it



The C version (slow version)

```
void merge_sort_slow(int* arr, size_t n) {  
    if (n <= 1) return;  
  
    int hi_start = ((n + 1) >> 1);  
    merge_sort_slow(arr, hi_start);  
    merge_sort_slow(arr + hi_start, n - hi_start);  
  
    // the slow part:  
    int* merge_buf = malloc(n * sizeof(int));  
    merge(arr, arr + hi_start, hi_start, n - hi_start, merge_buf);  
    memcpy(arr, merge_buf, sizeof(int) * n);  
    free(merge_buf);  
}
```

It's slow because we allocate repeatedly. [what tells us that allocation is happening?]
[Can you come up with a way to avoid this?]

Proving it

We showed that `merge` works. How do we show that `merge_sort` works?

We're going to have trouble using regular (weak) induction:

- The base case is fine: `mergesort([], 0)` is sorted.

- The problem is the inductive case:

suppose `mergesort(arr, n)` works for some `n` and `arr`

can we prove `mergesort([a] ++ arr, n + 1)` works?

The issue is that we split our array in half. So we have:

- `mergesort([a] ++ lo, (n + 2) / 2)`
- `mergesort(hi, n / 2)`

But our inductive hypothesis only works for `mergesort(arr, n)`. Not `lo` or `hi`.

The problem with weak induction

Weak induction assumes that we can go from $P(n)$ to $P(n + 1)$.

That means, in this case we assume that something about `mergesort(arr, n)` lets us prove `mergesort([a] ++ arr, n + 1)`.

But this isn't how mergesort works! It's not like insertion sort where it uses the fact that the array was sorted up to `i`, to generate a sorted array up to `i + 1`.

Instead, it breaks the array in half. We end up needing `mergesort([a] ++ half_sized_arr, (n + 2) / 2)` as an induction hypothesis, but we don't have one!

Weak induction principle

Remember that weak induction for natural numbers worked like this: We ask the induction principle algorithm for a proof of $P(2)$, where P is a proposition about a natural number. It recursively computes $P(1)$.

The proof of $P(1)$ recursively gets $P(0)$, which we provided, and then it applies $P(n) \implies P(n + 1)$ to generate $P(1)$, which it returns.

Then the $P(2)$ call takes $P(1)$ and applies $P(n) \implies P(n + 1)$ to it, yielding a proof of $P(2)$.

The issue is that the machine assumes that each proposition only depends on the one before.

Strong induction principle

But, what if we wanted to prove $P(10)$ using $P(5)$?

There's really no reason why we shouldn't be able to do this. If $P(2)$ and $P(3)$ depend on $P(1)$, and $P(4)$ depends on $P(2)$, and $P(5)$ depends on $P(3)$, then by the time we reach a high number, the proposition has been proven for the lower number it needs.

If we could use $P(5)$ in the proof of $P(10)$, then we could prove that mergesort works for array size 10, because it works for 5. How do we know it works for 5?

Because we proved it worked for 3 and 2?

How do we know that? Because we proved that it worked for 1.

The only wrinkle here is we had to prove it for 0 and also 1, because we cannot derive a proof of 1 from a proof of 0 in this case.

Strong induction principle (2)

The only way that strong induction differs from weak induction is in the inductive case:

$$\forall n \in \mathbb{N}, (\forall i \leq n, P(i)) \implies P(n+1)$$

Compare this to weak induction:

$$\forall n \in \mathbb{N}, P(n) \implies P(n+1)$$

Weak induction requires: "if we have proven the proposition for some number n , we can prove it for $n+1$ "

Strong induction requires: "if we have proven the proposition for *every number up to and including* n , we can prove it for $n+1$."

Let's demonstrate

Let our proposition be that for natural numbers n and arrays `arr`, where n is the length of `arr`, `mergesort(arr, n)` leaves `arr` sorted.

Does `mergesort([], 0)` work? Yes, it does. This is the base case.

There's a second base-case here: `mergesort({a}, 1)`. This also works.

Now, here's the key. We need to show:

$$(\forall i \leq n, P(i)) \implies P(n + 1)$$

So we "suppose" that `mergesort(arr', i)` works, for all up to and including `mergesort(arr, n)`, and we need to show that `mergesort({a} ++ arr, n + 1)` works.

Proving mergesort correct

Mergesort on n always breaks its array down into two pieces:

- One of length $n/2$ or $n/2 + 1$
- One of length $n/2$

Because these numbers are smaller than n , we can assume that we have proven that mergesort works for $n/2$ and $n/2 + 1$

Assuming that mergesort works, we end up with two sorted arrays, `lo` and `hi`. We merge them together, and we have already proven that `merge` works. Therefore, the result is a sorted list for any `n`. $\square$

How it works from 0 to n

One more time, consider this:

$P(0)$ and $P(1)$ were base cases. We know those work.

$P(2)$ only needs $P(1)$ to be proven. If mergesort works for $n = 1$, then we know it works for $n = 2$, because it will be split into two arrays of $n = 1$, mergesort will work correctly for them (we proved it), and then merge will work (we proved it, too).

$P(3)$ follows from $P(2)$ and $P(1)$, both of which are now proved.

$P(4)$ follows from $P(2)$, which followed from $P(1)$.

$P(5)$ follows from $P(3)$ and $P(2)$

Notice how high values of n always follow from lower values of n . This is why this induction principle is reasonable.

Practice

- Prove that $(n + 1) / 2$ gives us the first element of the second half. Use case analysis: case 1: $\text{len(lo)} = \text{len(hi)}$, case 2: $\text{len(lo)} = \text{len(hi)} + 1$. Answer is in `mod4.c` commented in the code for `merge_sort_slow` .
- Implement a faster merge sort that takes a pointer to a buffer and uses that instead of allocating.
- Write a binary search procedure. Given a sorted list, it will determine whether i is in the list by binary searching. Prove that it is correct using strong induction.
(If you don't prove it, and you haven't done binary search in a while, you *very likely* have a bug. Create a 100 element array of the numbers 1 to 100, and make sure every one of them is found in 7 checks or fewer)

Questions?

How fast is it?

What is mergesort's recurrence relation?

[Can we guess it?]

Recurrence relation of mergesort

$$T(0) = 1$$

$$T(1) = 1$$

$$T(n) = 2T(n/2) + \Theta(n)$$

The $\Theta(n)$ is from merge.

Why is merge $\Theta(n)$? Because it has to copy one element for every value in a sub-array.

(I'm simplifying somewhat. Sometimes it's not exactly cut in half, and one piece is one element longer than the other. Ignoring this will not change the big- Θ)

footnote:

(If we wanted to be super rigorous, we'd use these inductive cases:

$$T(2n) = 2T(n) + \Theta(2n), T(2n + 1) = T(n + 1) + T(n) + \Theta(2n + 1).$$

This way we handle the fact that odd lists don't break exactly evenly.

The big- Θ we get will be the same.)

How does this scale?

Every time we split in half, we have to solve 2 subproblems.

Once we're done with a sub-problem, we have to merge it back together, which takes $\Theta(n)$ time.

The reason this is hard to analyze is that the size of n keeps changing. When we're merging 2-element lists, it's very different than when we merge 1024-element lists.

So how can we analyze this?

Every stage

Let's express each level of recursion as a row:

[4, 2, 7, 8, 1, 3, 6, 5] split in half

[4, 2, 7, 8][1, 3, 6, 5] split each half in half

[4, 2][7, 8][1, 3][6, 5] split each quarter in half

[4][2][7][8][1][3][6][5] we're done splitting. This is "conquer"

[2, 4][7, 8][1, 3][5, 6] now merge each 8th

[2, 4, 7, 8][1, 3, 5, 6] now we merge the quarters

[1, 2, 3, 4, 5, 6, 7, 8] merge the halves

Splitting is "divide". Conquering here is just returning [a]. Recombine is merge.

Question: how many rows will there be for each phase?

Analyzing the number of rows

We want to count how many steps it takes to go from 1 array of n elements to n arrays of 1 element. That will be roughly half the number of rows.

For the first part: we're asking "how many times can we halve n until it hits 1."

This is the same as asking "how many times do we have to double 1 until it hits n "

This is the same as asking $2^k = n$, solve for k

[What is this?]

Analyzing the number of rows (2)

That's right, $\lg n = k$

So for $n = 8$, $\lg 8 = 3$. There will be three transformations: $8 \rightarrow 4$, $4 \rightarrow 2$, and $2 \rightarrow 1$.

This is the divide phase. We do 2 recursive calls, then 4, then 8, for a total of 14.

Once we hit 8 arrays of 1, we go the other way:

$$2 \times 1 = 2, 2 \times 2 = 4, 2 \times 4 = 8$$

This is the recombine phase. We write a byte for each element. 3 rows, 8 elements each.

Quick question: which one takes more $\Theta(1)$ operations?

Divide vs recombine

In this case, recombining is slower, but both phases are $\Theta(n \lg n)$.

In short: doubling the number of rows will mean doubling the number of splits, and also doubling the number of merged integers. So they grow at the same rate.

In some D&C algorithms, dividing takes more work. In others, recombining is slower. It turns out we typically get logarithms as the big- Θ when they're about the same speed.

Now that we've determined that mergesort is likely $\Theta(n \lg n)$, let's prove it. But first...

Practice

- Do the previous merge-sort analysis on $n = 4$ and $n = 16$. That is, sketch the number of elements visually, and then verify that dividing takes $\lg n$ rows and recombining takes $\lg n$ rows
- Are there any values for n where that formula doesn't work?

Questions

Proving this

The purpose of drawing a sketch of how many rows and columns there are was to give us a proposition for induction.

Here is the recurrence relation:

$$T(0) = 1$$

$$T(1) = 1$$

$$T(n) = 2T(n/2) + \Theta(n)$$

The proposition we want to prove is: $T(n) = \Theta(n \lg n)$

Let's start with $T(n) = O(n \lg n)$. I'll leave $T(n) = \Omega(n \lg n)$ to you.

Using induction

We can prove this using strong induction. But we don't know what values to use for n_0 and C . Let's just pretend we chose them so we can solve for them.

We want to show:

$$\forall n \geq n_0, T(n) \leq C(n \lg n)$$

We need to use induction, so let's write our subgoals:

$$0 \geq n_0 \implies T(0) = 1 \leq C(0 \lg 0)$$

$$\forall n, (\forall i < n, i \geq n_0 \implies T(i) \leq C(i \lg i)) \implies n \geq n_0 \implies T(n) \leq C(n \lg n)$$

Slight of hand warning: I modified the strong inductive principle. Instead of showing that if we prove $P(0)$ up to $P(n)$ that it implies $P(n + 1)$, I changed it to if we prove $P(0)$ up to $P(n - 1)$ it implies $P(n)$. This is equivalent, and it makes the math nicer for mergesort.

The base case

$$0 \geq n_0 \implies 1 \leq C(0 \lg 0)$$

Is the conclusion true? It's not true for $n_0 = 0$. Because $\lg 0$ is undefined.

If we choose $n_0 = 1$, it's vacuously true, because $0 < 1$.

We actually don't want vacuous truth here because it will complicate the inductive step.

If we pick $n_0 = 1$, then, when we want to (in the next step) show

$(\forall i \leq n, P(i)) \implies P(n+1)$, we actually won't be able to. $P(0) \implies P(1)$ is false with $n_0 = 1$, because $P(1)$ is $1 \leq C(1 \lg 1) = 0$. Basically there will be no C we can pick.

Therefore we *have* to choose $n_0 = 2$

As a rule, when using strong induction, we usually want to actually find the first non-vacuous case.

The (strong) inductive case

$$(\forall i < n, n \geq 2 \implies T(i) \leq C(i \lg i)) \implies n \geq 2 \implies T(n) \leq C(n \lg n)$$

We start by supposing this hypothesis: $(\forall i < n, i \geq 2 \implies T(i) \leq C(n \lg n))$

Then we suppose $n \geq 2$. We must show $T(n) \leq C(n \lg n)$

Let's simplify $T(n) = 2 \cdot T(\frac{n}{2}) + \Theta(n)$. We're assuming integer division.

Now, because of our inductive hypothesis, we can assume the proposition is true for $T(\frac{n}{2})$. That is: $T(\frac{n}{2}) \leq C \frac{n}{2} \lg \frac{n}{2}$

Now, every time we see $T(\frac{n}{2})$ we can replace it with $C \frac{n}{2} \lg \frac{n}{2}$ and get something bigger. We can use this to build an inequality. This is called the substitution method.

Using the substitution method

Start with $T(\frac{n}{2}) \leq C \cdot \frac{n}{2} \lg \frac{n}{2}$, this is the induction hypothesis.

$$\implies 2 \cdot T(\frac{n}{2}) \leq C \cdot n \lg \frac{n}{2}$$

$$\implies 2 \cdot T(\frac{n}{2}) \leq C \cdot n \lg \frac{n}{2} = C \cdot n(\lg n - \lg 2) = C \cdot n \lg n - n$$

So this shows:

$$2 \cdot T(\frac{n}{2}) + 1 + \Theta(n) \leq Cn \lg n = O(n \lg n) \quad \square$$

Questions

Practice

- We did most of the proof that mergesort is $O(n \lg n)$. Now complete the proof and show it is $\Omega(n \lg n)$.

Different kinds of D&C problem

Let's take a (brief) breather from sorting algorithms and consider something simpler.

Here is a binary search tree node structure:

```
typedef struct node_t {  
    struct node_t *l, *r;  
    void* data; // we don't care about this field for the exercise  
} Node;
```

Do you recall how to count the number of nodes in the tree?

Counting nodes

```
size_t count_nodes(Node* root) {  
    if (!root) return 0;  
    return count_nodes(root->l) + count_nodes(root->r) + 1;  
}
```

Prove it? By strong induction on the size of the tree n :


- When $n = 0$, the tree is empty and `count_nodes` returns 0.
- The size of the two subtrees are both smaller than the whole tree, so by the strong inductive hypothesis:

if `count_nodes(l)` and `count_nodes(r)` are correct, then `count_nodes(l) + count_nodes(r) + 1` adds the nodes in both sides of the tree, plus the root.

But what is the big- Θ ?

It doesn't really matter how the tree looks, but to keep the analysis simple, let's assume it's balanced and complete.

How does this look?



D&C analysis

During the divide phase, we have two recursive calls per inner node. There are 7 inner nodes, so that's 14 recursive calls.

For 7, there are 6 recursive calls. For 3 there are 2.

In general, there will always be $\lfloor n/2 \rfloor$ inner nodes in a balanced, complete tree, and therefore $2 \cdot \lfloor n/2 \rfloor$ function calls.

What about the recombine phase? It's just an addition. That's $\Theta(1)$

Recurrence relation

The recurrence relation for balanced trees is therefore:

$$T(0) = 1$$

$$T(n) = 2 \cdot T(n/2) + 1$$

Looks similar to the mergesort case, but that 1 instead of a $\Theta(n)$ changes things.

So what do you think? We have n -ish recursive calls, but recombining happens once per inner node.

$$\text{That's } \Theta(n) + \Theta(n) = \Theta(n)$$

Are you convinced that this is $\Theta(n)$?

Why is it different?

This time, the divide phase was more expensive than the recombine. Therefore we're basically counting the number of divisions.

For mergesort, they were about the same. We would divide proportional to n times, but then *each* recombine took $\Theta(n)$ for the number of elements in the recombine. The fact that the recombine got more expensive depending on how many times we had split was where the $\lg n$ factor came from.

I'm going to hold off on proving this one. You'll see why.

Questions

Practice

- Prove that if the tree is degenerate (i.e., every node has at most one child), counting is still $\Theta(n)$
- Draw a similar diagram for counting the depth of the tree instead of the number of nodes. Assume the tree is balanced, and then do it for a degenerate tree. Do these have the same big- Θ ?

One more example

Suppose we want to find the k^{th} largest value in a list of n items.

(This is a common tech-interview question. You can also efficiently solve it with an k -sized min-heap)

The naive approach is to just sort the list and take element `k - 1`. That is not optimal, however. Instead, we can improve on the time it takes to sort by recognizing that we don't need the whole list to be sorted. We just need to know the top k .

Think of it like this: if you want to know the max, it's clearly $\Theta(n)$. What if you want the 2^{nd} largest? It should be almost as fast.

One way to do it is with the quickselect algorithm.

Quickselect

Suppose the list has $n = 15$, and we want the 13th smallest value. (the 3rd largest)

Suppose we have the ability to magically determine the median value (we don't, but pretend we do).

We can split the list into three pieces:

$< \text{median} \quad ++ \quad \text{median} \quad ++ \quad \geq \text{median}$

That is, those elements that are less than the median, the median itself, and those that are greater.

Suppose the values are all different. We expect the lower half to have 7 elements, and the upper half to have the same.

Quickselect (2)

How does that help us?

Because now we know the 13th smallest value is in the upper list.

In fact, the 13th smallest of the whole list is the 5th smallest value in the upper list, because we know all 8 of the other values are not included in the upper list.

list: 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15

low: 1, 2, 3, 4, 5, 6, 7, mid: 8, high: 9, 10, 11, 12, 13, 14, 15

Notice that 13, which is the 13th smallest value in the whole thing, is actually number 5 in the high list.

(Note: obviously if the list is sorted we can easily find the 13th smallest value, but the algorithm I'm showing you works even if it's unsorted.)

Quickselect (3)

Now what? We recurse. We're looking for the 5th smallest value in...

list: 9, 10, 11, 12, 13, 14, 15

low: 9, 10, 11 , mid: 12 , high: 13, 14, 15

There are 3 values in the lower list, and 1 in the middle. We want the 5th, so we know that it is in the upper list. But because it's in the upper list and there are 4 smaller values, We don't want the 5th smallest, we want the 1st smallest in the upper list:

list: 13, 14, 15

low: 13 , mid: 14 , high: 15

Quickselect (4)

list: 13, 14, 15

low: 13 , mid: 14 , high: 15

Here, the midpoint gives us a number that's too large, so we know the solution is in the lower list. When the solution is in the lower list, we don't subtract the number of items from the one we're looking for.

list: 13

at this point, we want the 1st smallest value in 13 , which is 13 . We found it!

Partition

Of course, we don't magically have the ability to know the median. But we can get pretty close.

We can randomly choose a number in the array, and then put all the values smaller to the left, and all the values greater or equal to the right.

Sometimes the left array will be tiny, and sometimes the right array will be tiny.

However, we might get lucky and have the value we're looking for be in the tiny part. That would really speed things up.

From an information-theory point of view, ideally, the two halves will be the same size. This will minimize the expected number of searches.

Example partition function

Here's a simple partition in C, based on Tony Hoare's partition method:

```
size_t partition(int* arr, size_t n) {
    if (n <= 1) return 0;
    int r = arr[n - 1]; // choose the last element to be the pivot
    size_t i_lo = 0, i_hi = n - 2; // it's n - 2 because n - 1 is the pivot
    for (;;) {
        // find the first value bigger than the pivot, and the last smaller
        while (i_lo < n && arr[i_lo] < r) i_lo++; // could remove 1st check
        while (i_hi > 0 && arr[i_hi] >= r) i_hi--;

        if (i_lo < i_hi) swap(arr + i_lo, arr + i_hi);
        else break;
    }
    swap(arr + i_lo, arr + n - 1); // put the pivot into position
    return i_lo;
}
```

Explaining it

First, we select a "pivot", which is the value we want to be in the middle.

Actually figuring out the median would require us to sort the array, which is typically $\Theta(n \lg n)$. Not good.

Instead we just guess randomly by assuming the pivot is the last element.

Then, we enter a loop: `for(;;) ...`

Explaining it (2)

The first while loop searches for the first value that is smaller than the pivot:

```
while (i_lo < n && arr[i_lo] < r) i_lo++;
```

The second searches for the last value that is at least as big as the pivot:

```
while (i_hi > 0 && arr[i_hi] >= r) i_hi--;
```

Once we find both values, we swap them. This puts them in order.

```
if (i_lo < i_hi) swap(arr + i_lo, arr + i_hi);  
else break;
```

We want everything smaller than the pivot to the left.

Once `i_lo >= i_hi`, we're done. The small values are on the left.

Explaining it (3)

The last thing we do is

```
swap(arr + i_lo, arr + n - 1); // put the pivot into position  
return i_lo;
```

This makes it so that the pivot is actually greater than everything to its left and less than or equal to everything to its right.

We haven't sorted the list, but we've sorted the pivot.

Practice

- Prove that partition results in an array in which every element $\text{arr}[0..p) < \text{arr}[p]$ and $\text{arr}[p + 1..n) \geq \text{arr}[p]$.
- Implement quickselect using this partition. This is in `mod4.c`.
- Prove that it works.

Questions?

The big- Θ of quickselect average case

In the best case, quickselect finds the k th value instantly (it's the pivot). But, we had to partition it once to know that, so it's still $\Theta(n)$.

In the worst case, either the pivot keeps being the biggest value and we want the smallest, or vice versa. This means we pivot n times, and pivot is $\Theta(n)$, so it ends up being $\Theta(n^2)$ worse case.

But what is its average case? It's when there are an equal number of values on either side of the median.

Quickselect average case recurrence relation

$$T(0) = 1$$

$$T(n) = T(n/2) + \Theta(n)$$

Interesting. So we split it in half, but only once (it's not $2 \cdot T(n/2)$). Then we do something that takes roughly n time.

Sketch of quickselect

Imagine we have $n = 64$. We partition the 64, then identify which half has our value.

Then we partition the 32, then split roughly in half.

Then we partition 16...

8...4...2...1... and we found it.

In total, we have to process roughly $64 + 32 + 16 + 8 + 4 + 2 + 1$ values. Which is 127.

It's not a coincidence that this is $2n - 1$.

This algorithm ends up being $\Theta(n)$. Again, we'll prove it later.

Three algorithms

Mergesort had this recurrence:

$$T(0) = T(1) = 1$$

$$T(n) = 2 \cdot T(n/2) + \Theta(n)$$

count BST nodes had this recurrence for the balanced case:

$$T(0) = 1$$

$$T(n) = 2 \cdot T(n/2) + 1$$

and quickselect had this recurrence for the average case:

$$T(0) = 1$$

$$T(n) = T(n/2) + \Theta(n)$$

Notice any patterns?

The D&C pattern

Every divide and conquer algorithm has a recurrence relation that looks like this:

$$T(n) = A \cdot T(n/B) + f(n)$$


A is the number of splits we make

B is the size of each split. I.e., are we splitting in half, splitting in thirds, etc.

$f(n)$ is the function we do to recombine.

It turns out that knowing just A , B , and $f(n)$, we can totally characterize the big- Θ of the algorithm.

There's a theorem, called the "Master Theorem" that tells us how...



This relation looks like $3 \cdot T(n/3) + f(n)$

The question is, how much work is f doing?

It gets called for every inner node. That is, for every row except the last one. And each row down, it does less work. How many rows are there?

The number of rows

How many times can you split before reaching the leaves? In this case: $\log_3(n)$.

In general $\log_B(n)$

And how many nodes will there be on each row?

In this case, row r , starting at 0, will have 3^r elements.

In general: A^r nodes per row.

At row r , each node has $(n/3^r)$ elements. (n/B^r) in general.

So, in general, we invoke f like this:

$$f(n) + A(f(n/B)) + A^2(f(n/B^2)) + \dots = \sum_{i=0}^{(\log_B n)-1} (A^i f(n/B^i))$$

But that's not the only source of work we need to care about!

The number of leaves

When we reach a leaf, D&C algorithms end up being expressed as:

$T(0) = c$, or $T(1) = d$, or some kind of constant time base case. The key is that however many leaves there are, that's how many times we will call $T(1)$.

How many times does that happen?

This computation is easier: it's once per leaf. And there are $A^{\log_B n}$ leaves.

So we have a term: $A^{\log_B n}$.

Combining the two

Therefore, the total amount of work our algorithm will do is:

$$\sum_{i=0}^{(\log_B n)-1} (A^i f(n/B^i)) + A^{\log_B n}$$

for some constants A and B

This is a sum of two terms. So if one side dominates the other, then that's the big- Θ

How do we know if one side dominates the other?

(*) I'm using [Wikipedia](#) instead of the book for this definition.

Revisiting the sum

At level i of the recursion tree, there are A^i nodes.

Suppose there is a function g that approximates the amount of work at a node:
 $f(n) \approx g(n) = \Theta(n^c)$ for some c .

Then the cost per node is $\Theta(n/B^i)^c$. So the cost per row is $\Theta(n^c \cdot A^i / B^{ic})$.

Let $L = \log_B n$, how do we get $\sum_{i=0}^{L-1} \Theta(n^c \cdot A^i / B^{ic})$?

First, we can refactor a bit:

$$\Theta(n^c \cdot \sum_{i=0}^{L-1} A^i / B^{ic}) = \Theta(n^c \cdot \sum_{i=0}^{L-1} (A/B^c)^i)$$

Rebuilding $T(n)$

Now, let's add the other term back in to get our $T(n)$ expression:

$$T(n) = \Theta(n^c \cdot \sum_{i=0}^{L-1} (A/B^c)^i) + A^L$$

We're close. We want both terms to have a base of n . Then we can compare them.

We use the exponent/log swap trick:

$$A = B^{\log_B A}, A^L = A^{\log_B n} = (B^{\log_B A})^{\log_B n} = (B^{\log_B n})^{\log_B A}$$

And $B^{\log_B n}$ is just n , so: $A^L = n^{\log_B A}$ So now we have:

$$T(n) = \Theta(n^c \cdot \sum_{i=0}^{L-1} (A/B^c)^i) + \Theta(n^{\log_B A})$$

c_{crit}

$$T(n) = \Theta(n^c \cdot \sum_{i=0}^{(\log_B n)-1} (A/B^c)^i) + \Theta(n^{\log_B A})$$

The sum is actually a geometric series with $r = A/B^c$

So everything hinges on that ratio: $r = A/B^c$

(if you've forgotten: you learned about geometric series in one of your calculus classes)

One last substitution will give us something to make a decision on...

Let $c_{\text{crit}} = \log_B A$, then $A = B^{c_{\text{crit}}}$

So $r = A/B^c = B^{c_{\text{crit}}}/B^c = B^{c_{\text{crit}}-c}$

$$T(n) = \Theta(n^c \cdot \sum_{i=0}^{(\log_B n)-1} (B^{c_{\text{crit}}-c})^i) + \Theta(n^{c_{\text{crit}}})$$

Case 1: tree is leaf-heavy

If $A/B^c < 1$, that's equivalent to saying $B^{c_{\text{crit}}} < B^c \equiv B^{c_{\text{crit}}-c} < 1 \equiv C_{\text{crit}} > c$

If $C_{\text{crit}} > c$, $r < 1$, so the sum is going to result in a constant:

$$T(n) = \Theta(n^c \cdot \sum_{i=0}^{L-1} (A/B^c)^i) + \Theta(n^{c_{\text{crit}}}) = \Theta(n^c) + \Theta(n^{c_{\text{crit}}})$$

If $c_{\text{crit}} > c$, the whole thing simplifies to:

$$T(n) = \Theta(n^{c_{\text{crit}}})$$

That is, if $T(n) = A \cdot T(n/B) + O(n^c)$, and $\log A / \log B > c$, then $T(n) = \Theta(n^{c_{\text{crit}}})$

In other words: the work we do inside the tree is much smaller than the total number of leaves, so the number of leaves ends up being the big- Θ

Case 2: the tree and leaves are balanced

If $A/B^c = 1$, that's equivalent to saying $C_{\text{crit}} = c$

The thing inside the sum $= 1$, so we end up with a logarithm factor. Recall $L = \log_B n$

$$T(n) = \Theta(n^c \cdot \sum_{i=0}^{L-1} (A/B^c)^i) + \Theta(n^{c_{\text{crit}}}) = \Theta(n^c \cdot \log_B n) + \Theta(n^{c_{\text{crit}}}) = \Theta(n^c \cdot \log_B n)$$

We assumed that $f(n) = \Theta(n^c)$ to derive case 1, but if our original function had factors of $\log n$ in it, that is, if $f(n) = \Theta(n^c (\log n)^2)$ then we would end up with one more log power:

$$T(n) = \Theta(n^c (\log n)^2 \cdot \sum_{i=0}^{L-1} (A/B^c)^i) = \Theta(n^c (\log n)^2 \cdot \log n) = \Theta(n^c \cdot (\log n)^3)$$

In general, if $T(n) = A \cdot T(n/B) + \Theta(n^c (\log n)^k)$ and $c = \log A / \log B$,

$$T(n) = \Theta(n^c (\log n)^{k+1})$$

Note: most of the time, k will be 0. This will make $\log n$ magically appear.

Case 3: the weird one: inner work heavy

If $A/B^c > 1$, that's equivalent to saying $C_{\text{crit}} < c$

Here, the thing inside the sum is > 1 . This means that we're doing more work inside the tree than we are at the leaves. This is the least common case.

Recall that $A/B^c = B^{c_{\text{crit}}-c}$, so

$$\begin{aligned} T(n) &= \Theta(n^c \cdot \sum_{i=0}^{L-1} (A/B^c)^i) + \Theta(n^{c_{\text{crit}}}) \\ &= \Theta(n^c \cdot (1 + B^{c_{\text{crit}}-c} + (B^{c_{\text{crit}}-c})^2 + \dots)) + \Theta(n^{c_{\text{crit}}}) \end{aligned}$$

In that big- Θ , only the largest term actually matters. So it ends up being

$$= \Theta(n^c \cdot (B^{c_{\text{crit}}-c})^L) + \Theta(n^{c_{\text{crit}}}) = \Theta(n^c \cdot (B^{c_{\text{crit}}-c})^L), \text{ which was all an approximation of } f(n) \text{ (this was an early substitution we made).}$$

Case 3: the regularity condition

Therefore, if $T(n) = A \cdot T(n/B) + f(n)$, where $f(n) = \Omega(n^c)$ and $c > c_{\text{crit}}$,
 $T(n) = \Theta(f(n))$

Basically, the recombining is so expensive that it ends up dominating everything.

Except there's a special condition, called a regularity condition. There must be a $k < 1$:

$$A \cdot f\left(\frac{n}{B}\right) \leq k \cdot f(n)$$

Why? Because we want to ensure that the work at the root is bigger than the work at the other inner nodes. If it is, then the work at the root node dominates all the others. If it's not, then there's not an obvious bound, and we need to use another technique.

Case 3: what if the regularity condition is false?

Then we can't use the master theorem. Sorry.

This can happen with functions that get *bigger* when their input value gets smaller. So things like $f(n) = 1/n$ would fail this requirement.

Also, random-ish functions are like this. So if $f(n)$ were a random number or hash function in which n were not interpreted as a time but rather as a hash input, the total time can be anywhere from 0 to 2^{512} or something, but making n smaller does not make the result smaller.

The purpose of the regularity condition is to ensure that all the children together cost less than the parent. So even if it doesn't get strictly bigger, it can't "get smaller slower". Therefore, a logarithm would also fail. This is academic, because $f(n) = \log(n)$ cannot ever be case 3, because $n^{c_{crit}}$ can't be bounded by a log for any positive c_{crit} .

Regularity condition promise

Okay, I feel your pain. I won't ask you an example in which the regularity condition *fails*.

Therefore, you should expect to find $k!$

I didn't use to give this promise, so old quizzes might have an instance where it fails.

To be honest, I couldn't come up with a non-goofy example in which it failed, and I didn't want to give goofy questions. This stuff is hard enough as it is.

The master method: in summary

Try to express your problem as:

$T(n) = A \cdot T(n/B) + f(n)$, where A and B are constant.

(If you can't, the master method does not apply)

Let $c_{\text{crit}} = \frac{\log A}{\log B}$

1. If $f(n) = O(n^c)$ and $c_{\text{crit}} > c$

Then $T(n) = \Theta(n^{c_{\text{crit}}})$


2. If $f(n) = \Theta(n^{c_{\text{crit}}} \cdot (\log n)^k)$ and $c_{\text{crit}} = c$

Then $T(n) = \Theta(n^c \cdot (\log n)^{k+1})$

3. If $f(n) = \Omega(n^c)$ and $c_{\text{crit}} < c$ and $\exists k, A \cdot f(n/B) \leq k \cdot f(n)$

Then $T(n) = \Theta(f(n))$

4. Otherwise, you can't use the master method



Questions

Revisiting `count_nodes`

Remember `count_nodes` ? When it's balanced:

$$T(n) = 2 \cdot T(n/2) + 1$$

$$c_{\text{crit}} = \log 2 / \log 2 = 1$$

$$f(n) = 1 = O(n^0)$$

$1 > 0$, so we're in case 1: $T(n) = \Theta(n^{c_{\text{crit}}}) = \Theta(n^1)$

That's it, we don't have to prove anything by induction.

Revisiting mergesort

Is mergesort *really* $\Theta(n \lg n)$?

$$T(0) = T(1) = 1$$

$$T(n) = 2 \cdot T(n/2) + \Theta(n)$$

$$c_{\text{crit}} = \log 2 / \log 2 = 1$$

$$f(n) = \Theta(n^1 \cdot (\log n)^0)$$

$$c = 1 = c_{\text{crit}}$$

$$\text{Case 2: } T(n) = \Theta(n(\log n)^1) = \Theta(n \lg n)$$

Yup, it is. No messy induction or substitution.

Revisiting quickselect

What on earth is quickselect in the average case?

$$T(0) = 1$$

$$T(n) = 1 \cdot T(n/2) + \Theta(n)$$

$$c_{\text{crit}} = \log 1 / \log 2 = 0$$

$$f(n) = \Theta(n) = \Omega(n^1)$$

$$c = 1, c_{\text{crit}} = 0, c_{\text{crit}} < c$$

$\exists k < 1, A \cdot n/B \leq k \cdot n$? $n/2 \leq kn$, choose $k = 1/2$, it holds.

We're in case 3: $T(n) = \Theta(n)$

Practice

Quicksort is just like quickselect, but after partitioning we just quicksort each side.

- Implement quicksort
- What is the big- Θ of the best case, where the partition method perfectly splits the two arrays?
- In the worst case , the partition always picks the worst value (i.e., it's bigger or smaller than all the others).
 - What recurrence relation do we get then?
 - Is it eligible for the master method?
 - Find its big- Θ either way.

Appendix A: Practice Quiz

I'm going to generate 4 recurrence relations. Your job will be to use the master theorem to find their big- Θ if applicable, or say "cannot use master method if not". For case 3, you will need to show that the regularity condition applies.

For c_{crit} , you can use whatever logarithm base you want, as long as you use it for the numerator and denominator.

For example: for $T(n) = 3 \cdot T(n/9) + n$, $c_{\text{crit}} = \log_3 3 / \log_3 9 = 1/2$ and $c = 1$ (because $f(n) = n = n^1$)

Remember that finding c can be a little tricky. If the function is a polynomial like n^3 , then c is 3. However, if it's something like $|\sin n|$ then we need to realize that $|\sin n| = \Theta(1) = \Theta(n^0)$ (because it's always in the range $[0, 1]$). This lets us discover that $c = 0$.

Warning!

You will not be allowed to use a calculator in the quizzes or exams! I will give you logarithms that you can easily compute in your head.

I strongly recommend specifically drilling logarithms if you are struggling with this. Here's a good set of flashcards to make:

1. For one side, generate two natural numbers, A and B , which are both exponents of the same base, b . For example, 125 and 25, or 27 and 81
2. For the reverse side, write the result of $\log_b A$, $\log_b B$, and $c_{crit} = \frac{\log_b A}{\log_b B}$
3. Drill yourself on your ability to, looking at just A and B , compute their common base and ratio of logarithms.

I'm saying this because I gave a quiz with $T(n) = 81 \cdot T(n/3)$ and people kept asking to use their calculator.

Other log situations

I can also ask for logs where the exact value doesn't matter.

For example, $T(n) = 5 * T(n/2) + n^4$

I don't know what the $\log 5 / \log 2$ is, but I do know it's somewhere between $\log 4 / \log 2 = 2$ and $\log 8 / \log 2 = 3$.

So therefore, it must be smaller than 4, and this is a case 3 recurrence relation!

I recommend some flashcards with this kind of scenario, too.

Since we're CS majors, it's a good idea to learn to use [tools like this one](#).

Another warning!

You will have a 15 minute time limit for the quiz!

That means about 3.5 minutes per problem.

If you've been drilling these, I think it is reasonable, but you have to practice under a time constraint.

Appendix A: some samples

1. $T(n) = 2 \cdot T(n/2) + n$

2. $T(n) = T(n/3) + n^2$

3. $T(n) = n + 4 \cdot T(n/2)$

4. $T(n) = 256 \cdot T(n/2) + n^3 \cdot |\sin n|$

5. $T(n) = 49 \cdot T(n/7) + \Theta(n^2 \log n)$

6. $T(n) = 100 \cdot T(n/10) + 1$

7. $T(n) = T(n - 1) + \Theta(n)$

Appendix A: the answers

1. $c_{\text{crit}} = 1$, $f(n) = n^1 \cdot (\log n)^0$, so case 2. $T(n) = \Theta(n \lg n)$
2. $c_{\text{crit}} = 0$, $f(n) = \Omega(n^2)$, $c_{\text{crit}} < 2$ case 3. $\exists k < 1, n^2/2 \leq kn^2$?
Yes, case 3: $T(n) = \Theta(n^2)$
3. I flipped the terms around but it doesn't matter: $c_{\text{crit}} = \lg 4 / \lg 2 = 2$,
 $f(n) = O(n^1)$, $2 > 1$, so case 1. $T(n) = \Theta(n^{c_{\text{crit}}}) = \Theta(n^2)$
4. $c_{\text{crit}} = 7$, $f(n) = O(n^3)$, therefore case 1: $T(n) = \Theta(n^7)$
5. $c_{\text{crit}} = \log_7 49 / \log_7 7 = 2$, $c = 2$, case 2. $T(n) = \Theta(n^2 (\log n)^2)$
6. $c_{\text{crit}} = \log_{10} 100 / \log_{10} 10 = 2$, $1 = O(n^0)$, $1 > 0$, $c_{\text{crit}} > c$ case 1:
 $T(n) = \Theta(n^2)$
7. Not a divide and conquer problem; can't use the master theorem.

I'm confused about finding C!

One of the challenging parts of the master method is that sometimes c isn't there.

In this situation, it is: $T(n) = 2 \cdot T(n/4) + n^2$

We use $A = 2$, $B = 4$, $c_{\text{crit}} = 1/2$, and $c = 2$. This is a case 3.

But what about this? $T(n) = 2 \cdot T(n/4) + n|\sin n| + n$

Then we need to ask:

- Is this case 1? If so, we would need $n|\sin n| + n = O(n^c)$ and have $c < 1/2$. This isn't possible. $|\sin n|$ is as big as 1. Therefore, the tightest bound of $n|\sin n| + n$ is $O(n)$, which implies $c = 1$.
- Is this case 2? No for the same reason.
- Therefore, it must be case 3. Now we need to find a c so $n|\sin n| + n = \Omega(n^c)$. We can do this with $c = 1$. Regularity is easy to show with $f(n) = n^1$

Finding C (2)

So the flowchart still works, but sometimes we need to consider all 3 possibilities.

These are the more challenging problems.

Here are some more problems, unworked, some of which are more challenging.

Remember, 15 minutes for 4 problems!

Appendix A: unworked samples

1. $T(n) = 3T(n/2) + n^2$

2. $T(n) = T(n/4) + n \log n$

3. $T(n) = 5T(n/3) + n$

(the answer can have a log ratio in it when it's not something computable by hand)

4. $T(n) = 2T(n/2) + n^2 \log_8 n$

5. $T(n) = 9T(n/3) + n^2$

6. $T(n) = 2T(n/2) + n(\log n)^2$

7. $T(n) = T(n-1) + T(n-2) + 1$

Appendix A: tricky samples that I can still ask

Ponder these. All of them have definite solutions, and are solvable using math you know. I consider these to be reasonable exam problems, but you'll want to study them.

- $T(n) = T(n/2) + 2^n$

Hint: This one *does* pass the regularity test.

- $T(n) = T(n!)$

Hint: this is the factorial function

- $T(n) = 4 \cdot T(n/2) + n^2 + n$

- $T(n) = 8 \cdot T(n/2) + (\log n)^3$

Hint: $(\log n)^k = O(n)$ for all constant k

Questions?

Appendix B: Old fashioned sorts

In the 1890's, [Herman Hollerith](#) invented the card sorter.

You would record data records on a punch-card, containing, e.g., census information.

[The cards had a grid that could have holes punched](#). For example, if you were tabulating peoples' ages, you could punch a "3" in one row and a "2" in another to record "32", along with other data.

These cards could then be sorted into physical buckets, where a card counter could tabulate them. So you could know the number of people who were 32.

Multiple passes

The operator would place the cards in hoppers and run the machine.

The machine would sort one column into buckets. Say, the leading digit.

The operator could then take 1 bucket, e.g., the 3 bucket, set up 10 empty buckets, set the machine to sort on the next column, and dump the 3 bucket back into the hopper.

Now the new buckets would have all the 30s, 31s, 32s, etc.

These could each be tabulated by a card counter.

What is the big- Θ ?

Assume that we can sort into 10 buckets.

If we have 10 values to sort, we sort them in one pass.

If we have 100 values to sort, we sort them in $1 + 10 \times 1 = 11$ passes.
(i.e., 1 pass to give 10 buckets of 10 each. For each one, 1 pass)

If we have 1000, we sort them in $1 + 10 \times (1 + 10 \times 1) = 111$

If we have 10000, we sort them in $1 + 10 \times (1 + 10 \times (1 + 10)) = 1111$

A little surprising

Write the recurrence relation:

$$T(\leq 10) = 1$$

$$T(n) = 10 \cdot T(1/10) + 1$$

Applying the master theorem, this is $\Theta(n \lg n)$ But the fan-out is so wide, it feels almost linear!

Radix sort

This is called radix sort. What's cool is we can pick the radix:

- Base-2 means we sort the numbers based on a binary digit.
- Base-4 means we sort the numbers based on 2 binary digits grouped together.
- Base-16 means we sort on hex digits. a 32-bit number would be sorted in 8 passes.

Radix sort downsides

- If we sort from biggest to smallest digit (big-endian radix sort), then most datasets will have lots of 0 digits. How often does a number > 1 billion arise naturally? You end up wasting several passes
- We can fix this by sorting from smallest to largest (little-endian radix sort), and then we can stop sorting numbers as soon as they are less than the $\text{radix}^{\text{(number of sorts)}}$. But then we lose the nice property that sorting part way makes the numbers "almost sorted".
- We have to allocate buckets, which is slow.

Practice

- Implement radix sort. Make the radix configurable.
- What happens to the recurrence relation if we change bases? I.e., instead of using base-10, we use base-16?
- Does that change the big- Θ ?
- Radix sort seems to have a nicer recurrence relation than quicksort. Why do you think we don't often use it?
- Does radix sort require allocation? Can you think of a way to optimize base-2 to not require more than 1 additional buffer?

Mermaid source (1)

```
---
config:
  theme: redux
---
flowchart TD
  A["mergesort([5,3,1,2,4])"]
  A --> AL["mergesort([5,3,1])"]
  AL --> ALL["mergesort([5, 3])"]
  ALL --> ALLL["[5]"]
  ALL --> ALLH["[3]"]
  AL --> ALH[1]
  A --> AH["mergesort([2,4])"]
  AH --> AHL["[2]"]
  AH --> AHH["[4]"]

  ALLL & ALLH --> B35("merge([5], [3]) = (3, 5)")
  B35 & ALH --> B135("merge([3,5], [1]) = [1,3,5]")
  AHL & AHH --> B24("merge([2], [4]) = [2, 4]")

  B135 & B24 --> RESULT("merge([1, 3, 5] [2, 4]) = [1, 2, 3, 4, 5]")
  -->
```